

Current GPT research-run accounting

GPT-6 Astra Ultra, 26–27 September 2026. Twenty-six published results across five datasets. Scratch: from scratch. Tools: tools allowed.

Dataset	Setting	Variant	Code lines	Input [M]	Cached [M]	Output [M]	Elapsed [h:mm]	API est. [USD]
DBText	Scratch	A1	802	37.514	36.204	0.245	0:57	\$64.54
		A2	405	33.228	32.676	0.193	0:53	\$48.87
		A3	672	48.361	46.967	0.353	8:10	\$81.26
		A4	687	33.317	32.740	0.240	0:56	\$51.71
		A5	727	21.998	21.513	0.189	1:02	\$36.74
DBText	Tools	A1	599	28.624	27.971	0.149	0:56	–
		A2	349	43.797	42.940	0.194	1:01	–
		A3	326	25.399	24.872	0.118	0:33	–
		A4	711	68.184	66.862	0.311	1:48	–
		A5	669	36.817	35.817	0.212	1:21	–
Python	Scratch	A1	100	25.984	25.354	0.223	0:38	\$43.74
		A2	834	41.031	40.281	0.326	8:51	\$65.43
		A3	231	25.107	24.540	0.215	1:36	\$42.04
		A4	574	23.827	23.330	0.209	0:51	\$39.68
Yelp	Scratch	A1	132	13.699	13.326	0.135	0:23	\$24.48
		A2	602	19.749	19.274	0.196	0:35	\$34.68
		A3	157	24.428	23.906	0.159	0:30	\$38.13
		A4	386	27.369	26.795	0.227	0:33	\$44.93
OpenStack	Scratch	A1	157	22.822	22.298	0.168	0:26	\$36.82
		A2	870	13.720	13.211	0.171	8:15	\$27.91
		A3	1,008	11.037	10.716	0.108	0:17	\$20.03
		A4	541	28.198	27.610	0.216	0:49	\$45.46
SQLStorm TPC-DS	Scratch	A1	225	28.657	28.126	0.210	0:45	\$44.93
		A2	119	17.215	16.781	0.162	0:25	\$30.06
		A3	92	22.230	21.773	0.175	0:34	\$35.95
		A4	148	15.217	14.836	0.138	0:27	\$26.34
Total recorded API estimate								\$883.75

Cost. The 21 API results total \$883.746020 in recorded token-based estimates, including all model sessions, failed attempts, restarts and continuations. These are not invoices. One additional \$8.450125 reservation for DBText Scratch A3 is unresolved and is excluded. A dash denotes a ChatGPT subscription result with no attributable per-result dollar charge in the records; it does not mean zero cost. Separate supervisor work, subscription fees, taxes and machine costs are excluded.

Usage. Input includes cached input; output includes reasoning. M = one million tokens.

Code lines. Nonblank, noncomment lines in the published custom source, headers and fitting/build scripts, summed across files and counted once per file. Imported libraries, supplied interfaces, metadata and generated data-only files are excluded. Counts describe the selected solution, not abandoned candidates; they depend on source formatting. Per-file counts and exclusions accompany the table.

Elapsed. Wall-clock span across all model sessions, including tools and recovery pauses. API: first request to last response. Subscription: first native turn start to last token event. The long Python/OpenStack A2 and DBText A3 spans include recovery waits. This is not model compute time.

Source: saved campaign ledgers and deduplicated native token events. Exact counts, timestamps, result IDs and historical rates accompany the table in CSV and JSON. The older 22 September accounting table is separate.